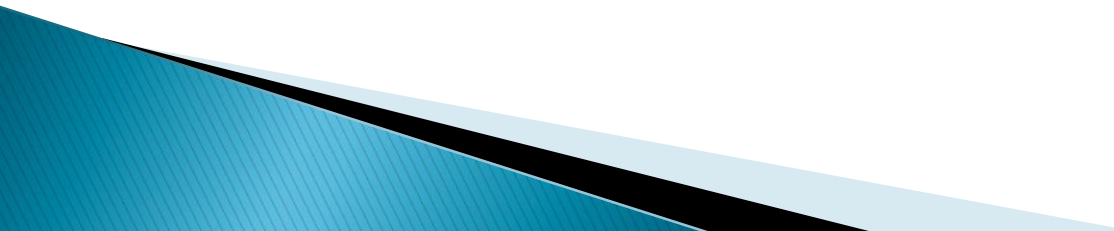
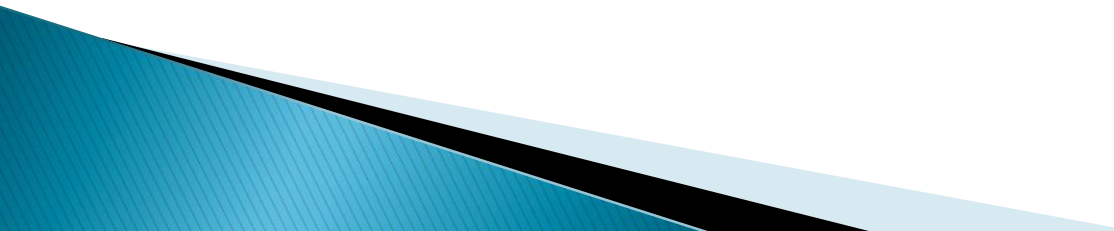


UNIT – II CLASSIFICATION

Unit II Classification

- ▶ Binary and Multiclass Classification:
Assessing Classification Performance,
Handling more than two classes, Multiclass
Classification–One vs One, One vs Rest Linear
Models: Perceptron, Support Vector Machines
(SVM), Soft Margin SVM, Kernel methods for
non-linearity
- 

Some Important Definitions

- ▶ Instance
 - Objects of Interests
 - ▶ Instance space
 - Set of all possible instances
 - ▶ Label Space
 - ▶ Output space
 - ▶ Model
 - Mapping from the instance space to output space
- 

Classification

- ▶ *A classifier is a mapping*



$$\hat{c} : X \rightarrow C$$

- ▶ *where $C = \{C_1, C_2, \dots, C_K\}$ is a finite and usually small set of class labels*

Classification

- ▶ Binary
- ▶ Multiclass

Binary Classification

- ▶ Also called as concept Learning if positive class can be meaningfully called concept.
- ▶ Simplest classifier with two classes
 - Positive  or +1
 - Negative  or -1

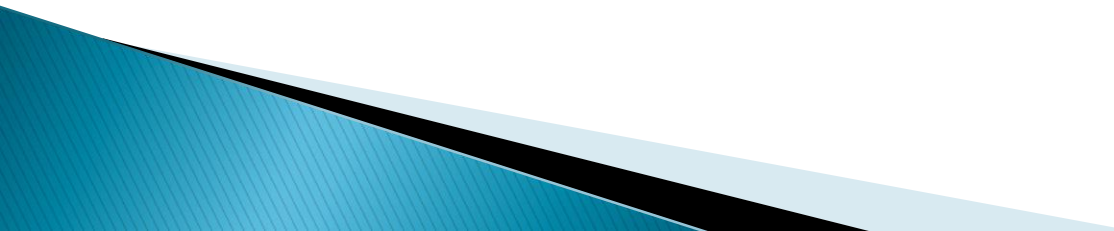
Performance of classification

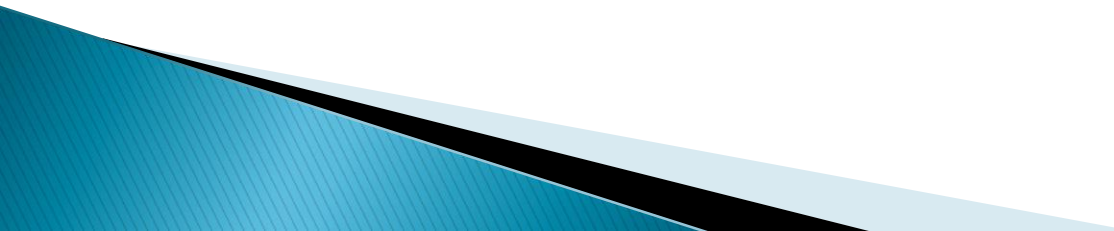
- ▶ Confusion Matrix or Contingency Table

Performance of classification

- Confusion Matrix or Contingency Table

	Predicted \oplus	Predicted \ominus	
Actual \oplus	30	20	50
Actual \ominus	10	40	50
	40	60	100

- ▶ True Positives
 - Correctly classified positives
 - ▶ True Negatives
 - Correctly classified negatives
 - ▶ False Negatives
 - Misclassified negatives
 - (incorrectly predicted as Negative)
 - ▶ False Positives
 - Misclassified positives
 - (incorrectly predicted as Positive)
- 

- ▶ True positive rate(tpr)
 - Proportion of positives correctly placed
 - Also called as sensitivity
 - ▶ True negative rate(tnr)
 - Proportion of negatives correctly placed
 - Also called as specificity
 - ▶ False positive rate(fpr)
 - ▶ False negative rate(fnr)
- 

Performance of classification

► Confusion Matrix or Contingency Table

	Predicted \oplus	Predicted \ominus	
Actual \oplus	30	20	50
Actual \ominus	10	40	50
	40	60	100

tpr = 60%

tnr = 80%

fnr = 40%

fpr = 20%



Precision and Recall

- ▶ Recall : sometimes called as sensitivity or specificity (TPR)
- ▶ Precision is a counterpart to true positive rate in the following sense:
 - precision is the proportion of actual positives among the predicted positives.





Classification



	Yes	No	
Yes	100	5	105
No	10	50	60
	110	55	165

165 patients were tested for the presence of the disease. Out of those 165 cases, the classifier predicted "yes" 110 times, and "no" 55 times. In reality, 105 patients in the sample have the disease, and 60 patients do not.

	Yes	No	
Yes	100	5	105
No	10	50	60
	110	55	165

Calculate

tpr, tnr, fpr, fnr, pos, neg, accuracy, error rate
precision and recall

- ▶ $\text{pos} = 105/165 = 0.64$
- ▶ $\text{neg} = 60/165 = .36$
- ▶ $\text{tpr} = 100/105 = 0.95$
- ▶ $\text{tnr} = 50/60 = 0.83$
- ▶ $\text{fpr} = 10/60 = 0.17$
- ▶ $\text{fnr} = 5/105 = 0.05$
- ▶ $\text{Accuracy} = (\text{TP} + \text{TN}) / \text{total} = (100 + 50) / 165 = 0.91$
- ▶ $\text{or acc} = \text{pos} \cdot \text{tpr} + \text{neg} \cdot \text{tnr} = 0.64 * 0.95 + 0.36 * 0.83 = 0.91$
- ▶ **Error rate**
 - $1 - \text{acc} = 0.09$
- ▶ **Or**
 - $(\text{FP} + \text{FN}) / \text{total} = (10 + 5) / 165 = 0.09$
- ▶ **Precision:**
 - $\text{TP} / \text{predicted yes} = 100 / 110 = 0.91$

Performance Measure

<i>Measure</i>	<i>Definition</i>	<i>Equal to</i>	<i>Estimates</i>
number of positives	$Pos = \sum_{x \in Te} I[c(x) = \oplus]$		
number of negatives	$Neg = \sum_{x \in Te} I[c(x) = \ominus]$	$ Te - Pos$	
number of true positives	$TP = \sum_{x \in Te} I[\hat{c}(x) = c(x) = \oplus]$		
number of true negatives	$TN = \sum_{x \in Te} I[\hat{c}(x) = c(x) = \ominus]$		
number of false positives	$FP = \sum_{x \in Te} I[\hat{c}(x) = \oplus, c(x) = \ominus]$	$Neg - TN$	
number of false negatives	$FN = \sum_{x \in Te} I[\hat{c}(x) = \ominus, c(x) = \oplus]$	$Pos - TP$	
proportion of positives	$pos = \frac{1}{ Te } \sum_{x \in Te} I[c(x) = \oplus]$	$Pos / Te $	$P(c(x) = \oplus)$
proportion of negatives	$neg = \frac{1}{ Te } \sum_{x \in Te} I[c(x) = \ominus]$	$1 - pos$	$P(c(x) = \ominus)$
class ratio	$clr = pos / neg$	Pos / Neg	
(*) accuracy	$acc = \frac{1}{ Te } \sum_{x \in Te} I[\hat{c}(x) = c(x)]$		$P(\hat{c}(x) = c(x))$
(*) error rate	$err = \frac{1}{ Te } \sum_{x \in Te} I[\hat{c}(x) \neq c(x)]$	$1 - acc$	$P(\hat{c}(x) \neq c(x))$

Performance Measure


Measure	Definition	Equal to	Estimates
true positive rate, sensitivity, recall	$tpr = \frac{\sum_{x \in Te} I[\hat{c}(x)=c(x)=\oplus]}{\sum_{x \in Te} I[c(x)=\oplus]}$	TP/Pos	$P(\hat{c}(x) = \oplus c(x) = \oplus)$
true negative rate, specificity	$tnr = \frac{\sum_{x \in Te} I[\hat{c}(x)=c(x)=\ominus]}{\sum_{x \in Te} I[c(x)=\ominus]}$	TN/Neg	$P(\hat{c}(x) = \ominus c(x) = \ominus)$
false positive rate, false alarm rate	$fpr = \frac{\sum_{x \in Te} I[\hat{c}(x)=\oplus, c(x)=\ominus]}{\sum_{x \in Te} I[c(x)=\ominus]}$	FP/Neg = $1 - tnr$	$P(\hat{c}(x) = \oplus c(x) = \ominus)$
false negative rate	$fnr = \frac{\sum_{x \in Te} I[\hat{c}(x)=\ominus, c(x)=\oplus]}{\sum_{x \in Te} I[c(x)=\oplus]}$	FN/Pos = $1 - tpr$	$P(\hat{c}(x) = \ominus c(x) = \oplus)$
precision, confi- dence	$prec = \frac{\sum_{x \in Te} I[\hat{c}(x)=c(x)=\oplus]}{\sum_{x \in Te} I[\hat{c}(x)=\oplus]}$	TP/(TP + FP)	$P(c(x) = \oplus \hat{c}(x) = \oplus)$



Coverage plot



- ▶ The number of positives *Pos*,
- ▶ *The number of negatives Neg, t*
- ▶ *The number of true positives TP*
- ▶ *The number of false positives FP.*
- ▶ A coverage plot visualises these four numbers by means of a rectangular coordinate system and a point

Coverage plot



- ▶ C1 has both more true positives and fewer false positives than C2, and so is better in both respects
- ▶ If one classifier outperforms another classifier on all classes, the first one is said to *dominate the second*

Coverage plot



- ▶ Consider a third classifier C3, better than C1 on the positives but worse on the negatives
- ▶ Although both C1 and C3 dominate C2, neither of them dominates the other. Which one we prefer depends on whether we put more emphasis on the positives or on the negatives.

Coverage plot

- ▶ Notice that the line segment connecting C1 and C3 has a slope of 1
- ▶ whenever we gain a true positive, we also lose a true negative (or gain a false positive)
- ▶ This doesn't affect the sum of true positives and true negatives, and hence the accuracy is the same wherever we are on the line.
- ▶ It follows that C1 and C3 have the same accuracy.
- ▶ In a coverage plot, classifiers with the same accuracy are connected by line segments with slope 1.

- ▶ consider the average of the true positive rate and the true negative rate, which we will call *average recall*, denoted *avg-rec*.

- ▶ On a line with slope 1 we have

$$\text{avg-rec} = (\text{tpr} + \text{tnr})/2 = (\text{tpr} + 1 - \text{fpr})/2 = (1 + y_0)/2,$$
which is a constant


- ▶ In a normalised coverage plot, line segments with slope 1 connect classifiers with the same average recall.

ROC Curve


- ▶ Receiver operating Characteristics plot
- ▶ popular measure for evaluating classifier performance
- ▶ based on two basic evaluation measures – specificity and sensitivity.
- ▶ It is a plot of TPR (sensitivity) against FPR(1 – specificity)
- ▶ The ROC plot uses FPR on the x-axis and TPR on the y-axis.

Threshold	Sensitivity	Specificity	1 - specificity
1	0.0	1.0	0.0
2	0.5	0.75	0.25
3	0.75	0.5	0.5
4	1.0	0.0	1.0

A ROC curve connecting 4 ROC points




Two ROC curves




- ▶ Area under the curve (AUC)
- ▶ If AUC is greater then that classifier is better
- ▶ If a classifier is perfect the AUC is 1
- ▶ If classifier makes all random guesses ,its AUC is 0.5


- ▶ The AUC score can be calculated by the trapezoidal rule, which is adding up all trapezoids under the curve. The areas of the three trapezoids 1, 2, 3 are 0.0625, 0.15625, and 0.4375. The AUC score is then 0.65625.

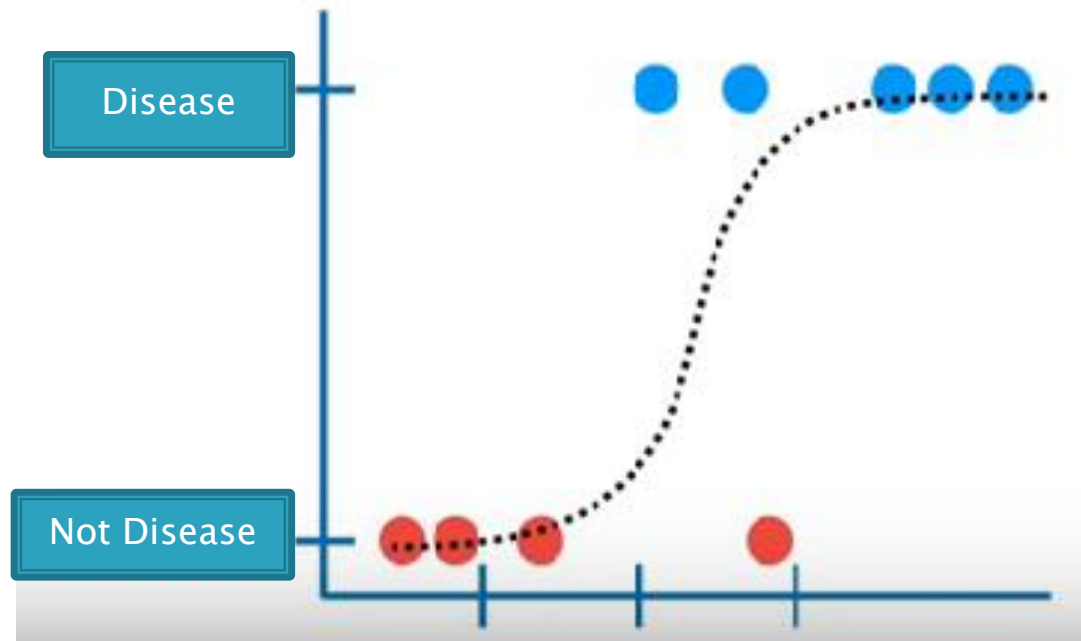



AUC scores




Let's start with some data...







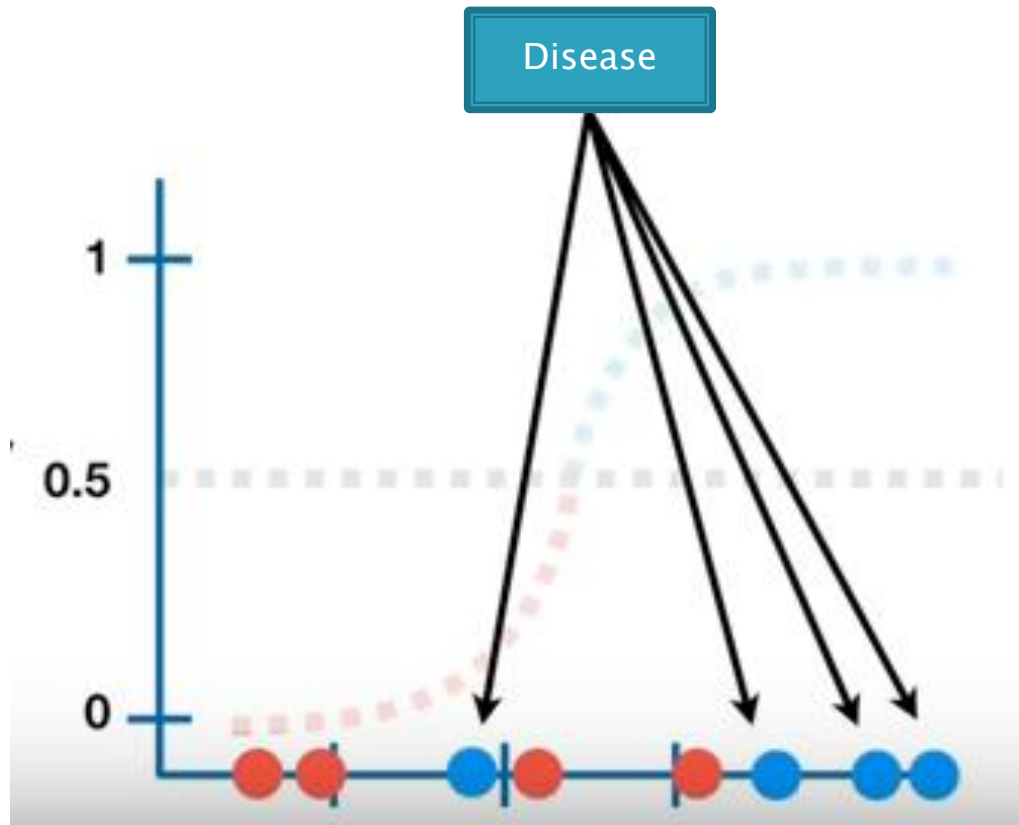





The
Probability of
Disease








Threshold 0.5




	Is having disease	Is not having disease
Is having disease	3	1
Is not having disease	1	3

Threshold 0.1




	Is having disease	Is not having disease
Is having disease	4	0
Is not having disease	2	2

Threshold 0.9




	Is having disease	Is not having disease
Is having disease	3	1
Is not having disease	0	4


So instead of being overwhelmed
with confusion matrices,
Receiver Operator Characteristic
(ROC) graphs
provide a simple way to summarize
all of the information.



True Positive Rate = Sensitivity = $\frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$



$$\text{False Positive Rate} = (1 - \text{Specificity}) = \frac{\text{False Positives}}{\text{False Positives} + \text{True Negatives}}$$




	Is having disease	Is not having disease
Is having disease	4	0
Is not having	4	0

$Tpr = \text{sensitivity} = 1$
 $Fpr = 1 - \text{specificity} = 1$




	Is having disease	Is not having disease
Is having disease	4	0
Is not having	3	1


$Tpr = \text{sensitivity} = 1$
 $Fpr = 1 - \text{specificity} = 0.75$




	Is having disease	Is not having disease
Is having disease	4	0
Is not having	2	2

$Tpr = \text{sensitivity} = 1$
 $Fpr = 1 - \text{specificity} = 0.5$







The **ROC** graph summarizes all of the confusion matrices that each threshold produced.




The **AUC** (Area Under the Curve) is **0.9**




The **AUC** makes it easy to compare one **ROC** curve to another.




The **AUC** for the **red ROC** curve is greater than the **AUC** for the **blue ROC** curve, suggesting that the **red** curve is better.



- ▶ Other matrices can be used to draw ROC
- ▶ FPR can be replaced by Precision
- ▶ ROC curves make it easy to identify the best threshold for decision making

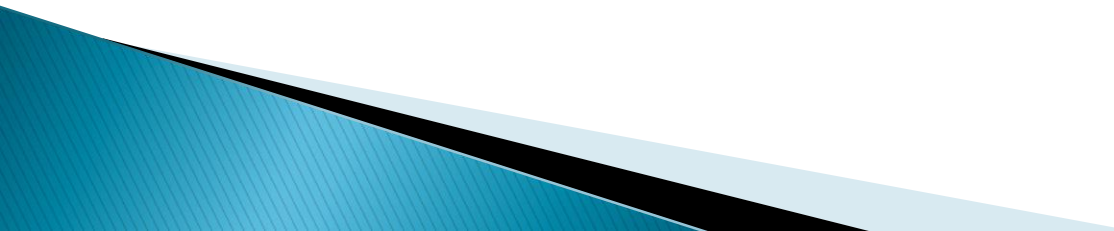


- ▶ And AUC can help you to choose which classification method is better





Multiclass Classification

- ▶ Classification task with more than two classifiers.
 - ▶ OVA
 - ▶ OVO
- 

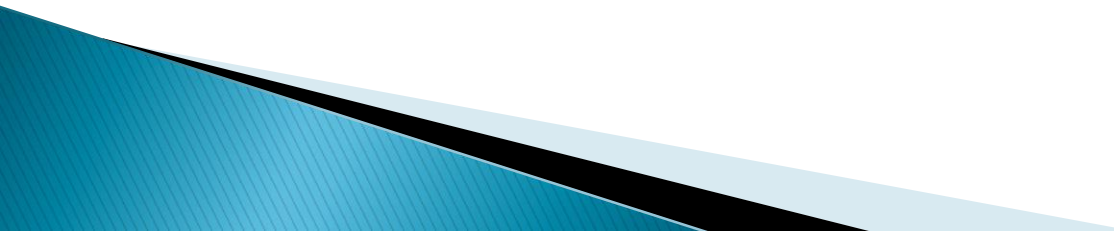
Multi-class classification

- ▶ One versus All(one-versus-rest)scheme
 - To train k binary classifiers, the first of which separates class C_1 from C_2, \dots, C_n , the second of which separates C_2 from all other classes, and so on.
 - To train k binary classifiers, the first of which separates class C_1 from C_2, \dots, C_n , the second of which separates *C_2 from all other classes, and so on.*

Multi-class classification

- ▶ one-versus-one scheme
 - we train $k(k - 1)/2$ binary classifiers, one for each pair of different classes.
 - If a binary classifier treats the classes asymmetrically, as happens with certain models, it makes more sense to train two classifiers for each pair, leading to a total of $k(k - 1)$ classifiers.

Output code matrix

- ▶ A convenient way to describe all these schemes to decompose a k -class task into l binary classification tasks is by means of output code matrix.
 - ▶ This is a k -by- l matrix whose entries are $+1$, 0 or -1 .
 - ▶ Each column of these matrices describes a binary classification task, using the class corresponding to the row with the $+1$ entry as positive class and the class with the -1 entry as the negative class.
- 

Output code matrix

- ▶ Output code matrix for one versus rest scheme

$$\begin{array}{l} \rightarrow c_1 \\ c_2 \\ c_3 \end{array} \begin{pmatrix} +1 & -1 & -1 \\ -1 & +1 & -1 \\ -1 & -1 & +1 \end{pmatrix}$$

$$\begin{pmatrix} +1 & 0 \\ -1 & +1 \\ -1 & -1 \end{pmatrix}$$

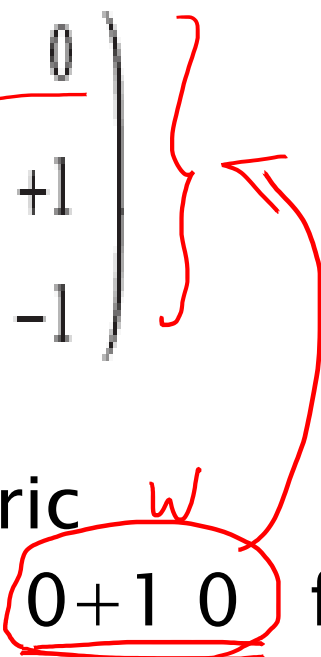
- ▶ unordered
- ▶ ordered
- ▶ Predict -1 +1 -1 for unordered scheme

Output code matrix

- ▶ Output code matrix for one versus one scheme

$C_1 \rightarrow$

$\begin{pmatrix} +1 & +1 & 0 \\ -1 & 0 & +1 \\ 0 & -1 & -1 \end{pmatrix}$	$\begin{pmatrix} +1 & -1 & +1 & -1 & 0 & 0 \\ -1 & +1 & 0 & 0 & +1 & -1 \\ 0 & 0 & -1 & +1 & -1 & +1 \end{pmatrix}$
---	---



- ▶ Symmetric w asymmetric
- ▶ Predict 0+1 0 for symmetric scheme
- ▶ the nearest code word is the first row in the matrix, and so we should predict C_1 .

- ▶ By using the formula the distance between a word W and a code word c is calculated

✓

$$d(\underline{w}, \underline{c}) = \sum_i \underline{(1 - w_i c_i) / 2}$$

- ▶ Minimum distance is the class we predicted

$$\rightarrow \begin{pmatrix} c_1 & c_2 & c_3 \\ +1 & +1 & 0 \\ -1 & 0 & +1 \\ 0 & -1 & -1 \end{pmatrix}$$

$c_1 \quad c_2 \quad c_3$

$w_1 \quad w_2 \quad w_3$
 ▶ Predict 0 + 1 0

Hamming distance.

C_1

$$d(w, c) = \sum_{i=1}^3 (1 - w_i c_i) / 2 \quad \checkmark$$

$$d(w, \underline{c_1}) = \frac{(1-0) + (1-1) + (1-0)}{2} = \underline{1} \quad \checkmark$$

$$d(w, \underline{c_2}) = 1.5$$

$$d(w, \underline{c_3}) = 1.5$$

3 class.

	Predicted				
Actual	15	2	3	20	A
	7	15	8	30	B
	2	3	45	50	C
	24	20	56	100	

$$\frac{15 + 15 + 45}{100} = \frac{75}{100}$$

$$0.75$$

per class precision A

$$\frac{15}{24} = 0.63$$

Recall A

$$\frac{15}{20} = 0.75$$

$$\text{precision B} = \frac{15}{20} = 0.75$$

recall B

$$\text{pre C} = \frac{45}{50} = 0.9$$

rec C


$$\left\{ \begin{array}{l} \text{Weighted Avg.} \\ \text{prec.} \end{array} \right. = 0.63 * 0.2 + 0.75 * 0.3 + 0.9 * 0.5 = 0.75$$

Accuracy as a weighted average

	Predicted \oplus	Predicted \ominus	
Actual \oplus	60	15	75
Actual \ominus	10	15	25
	70	30	100

From this table, we see that the true positive rate is $tpr = 60/75 = 0.80$ and the true negative rate is $tnr = 15/25 = 0.60$. The overall accuracy is $acc = (60 + 15)/100 = 0.75$, which is no longer the average of true positive and negative rates. However, taking into account the proportion of positives $pos = 0.75$ and the proportion of negatives $neg = 1 - pos = 0.25$, we see that

$$acc = pos \cdot tpr + neg \cdot tnr$$



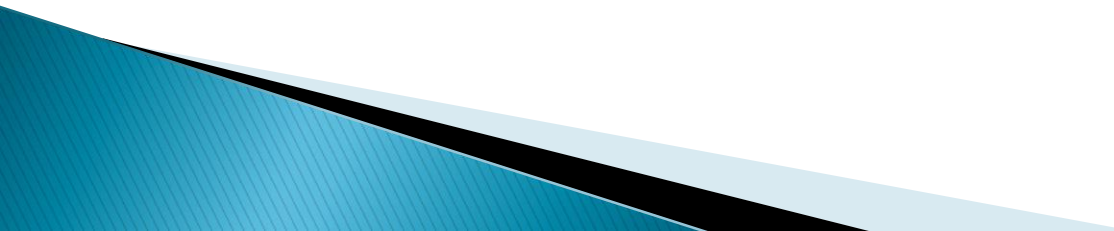
1) Linear Seperable

o o o o	* * * *
---------	---------





Perceptron

- ▶ Linearly Separable
 - ▶ Perceptron : Linear Classifier that will achieve perfect separation on Linearly Separable data
 - ▶ Simple NN
 - ▶ Learning rate(η)
- 

- ▶ Algorithm Perceptron(D, η)
 - i/p: Training Data and learning rate
 - o/p : weight vector w
- ▶ $w \leftarrow 0$
- ▶ converged \leftarrow false
- ▶ while converged = false do
 - Converged \leftarrow true
 - for($i=1$ to $|D|$)
 - If($y_i w \cdot x_i \leq 0$)
 - $W \leftarrow w + \eta y_i x_i$
 - Converged \leftarrow false
 - end
 - end
- ▶ end

- ▶ Every time an example is misclassified we add $y_i x_i$ to weight vector.
- ▶ Weight vector can be expressed as

$$w = \sum_{i=1}^n \alpha_i y_i x_i$$

Dual Perceptron

Algorithm 7.2: *DualPerceptron(D)* – perceptron training in dual form.

Input : labelled training data D in homogeneous coordinates.

Output : coefficients α_i defining weight vector $w = \sum_{i=1}^{|D|} \alpha_i y_i \mathbf{x}_i$.

```
1  $\alpha_i \leftarrow 0$  for  $1 \leq i \leq |D|$ ;  
2 converged  $\leftarrow$  false;  
3 while converged = false do  
4   converged  $\leftarrow$  true;  
5   for  $i = 1$  to  $|D|$  do  
6     if  $y_i \sum_{j=1}^{|D|} \alpha_j y_j \mathbf{x}_i \cdot \mathbf{x}_j \leq 0$  then  
7        $\alpha_i \leftarrow \alpha_i + 1$ ;  
8       converged  $\leftarrow$  false;  
9     end  
10  end  
11 end
```

Algorithm 7.4: $\text{KernelPerceptron}(D, \kappa)$ – perceptron training algorithm using a kernel.

Input : labelled training data D in homogeneous coordinates;
kernel function κ .

Output : coefficients α_i defining non-linear decision boundary.

```
1  $\alpha_i \leftarrow 0$  for  $1 \leq i \leq |D|$ ;  
2  $\text{converged} \leftarrow \text{false}$ ;  
3 while  $\text{converged} = \text{false}$  do  
4    $\text{converged} \leftarrow \text{true}$ ;  
5   for  $i = 1$  to  $|D|$  do  
6     if  $y_i \sum_{j=1}^{|D|} \alpha_j y_j \kappa(\mathbf{x}_i, \mathbf{x}_j) \leq 0$  then  
7        $\alpha_i \leftarrow \alpha_i + 1$ ;  
8        $\text{converged} \leftarrow \text{false}$ ;  
9     end  
10  end  
11 end
```



Basic operation on vectors in \mathbb{R}^n

1. Multiplication by a scalar

Consider a vector $\vec{a} = (a_1, a_2, \dots, a_n)$ and a scalar c

Define: $c\vec{a} = (ca_1, ca_2, \dots, ca_n)$

When you multiply a vector by a scalar, you “stretch” it in the same or opposite direction depending on whether the scalar is positive or negative.




Basic operation on vectors in \mathbb{R}^n

2. Addition

Consider vectors $\vec{a} = (a_1, a_2, \dots, a_n)$ and $\vec{b} = (b_1, b_2, \dots, b_n)$

Define: $\vec{a} + \vec{b} = (a_1 + b_1, a_2 + b_2, \dots, a_n + b_n)$




Recall addition of forces in classical mechanics.

Basic operation on vectors in \mathbb{R}^n

3. Subtraction

Consider vectors $\vec{a} = (a_1, a_2, \dots, a_n)$ and $\vec{b} = (b_1, b_2, \dots, b_n)$

Define: $\vec{a} - \vec{b} = (a_1 - b_1, a_2 - b_2, \dots, a_n - b_n)$



What vector do we need to add to \vec{b} to get \vec{a} ? I.e., similar to subtraction of real numbers.


Basic operation on vectors in \mathbb{R}^n

4. Euclidian length or L2-norm

Consider a vector $\vec{a} = (\underline{a_1}, \underline{a_2}, \dots, \underline{a_n})$

Define the L2-norm: $\|\vec{a}\|_2 = \sqrt{a_1^2 + a_2^2 + \dots + a_n^2}$

We often denote the L2-norm without subscript, i.e. $\|\vec{a}\|$



L2-norm is a typical way to measure length of a vector; other methods to measure length also exist.



Basic operation on vectors in \mathbb{R}^n

5. Dot product

Consider vectors $\vec{a} = (a_1, a_2, \dots, a_n)$ and $\vec{b} = (b_1, b_2, \dots, b_n)$

Define dot product: $\vec{a} \cdot \vec{b} = a_1b_1 + a_2b_2 + \dots + a_nb_n = \sum_{i=1}^n a_ib_i$ ✓


The law of cosines says that $\vec{a} \cdot \vec{b} = \|\vec{a}\|_2 \|\vec{b}\|_2 \cos \theta$ where θ is the angle between \vec{a} and \vec{b} . Therefore, when the vectors are perpendicular $\vec{a} \cdot \vec{b} = 0$.





5. Dot product (continued)


$$\vec{a} \cdot \vec{b} = a_1 b_1 + a_2 b_2 + \dots + a_n b_n = \sum_{i=1}^n a_i b_i$$


- Property: $\vec{a} \cdot \vec{a} = a_1 a_1 + a_2 a_2 + \dots + a_n a_n = \|\vec{a}\|_2^2$
- In the classical regression equation $y = \vec{w} \cdot \vec{x} + b$
the response variable y is just a dot product of the
vector representing patient characteristics (\vec{x}) and
the regression weights vector (\vec{w}) which is common
across all patients plus an offset b .















Support Vectors


- ▶ The data points nearest to the hyperplane
- ▶ the points of a data set that
- ▶ if removed, would alter the position of the dividing hyperplane.
- ▶ They can be considered the critical elements of a data set.



SVM




Support Vector Machines

- ▶ supervised classification techniques
- ▶ **supervised binary classification**
- ▶ The goal of the SVM is to train a model that assigns new unseen objects into a particular category.
- ▶ The SVM is an effective tool in high-dimensional spaces
- ▶ **Memory Efficiency**






support vectors.



SVM

- ▶ The training examples nearest to decision boundary are called support vectors.
- ▶ Margin is the distance between the decision boundary and the nearest training instance measured along w

$$y = mx + c$$



- ▶ Margin = $2m / ||W||$
- ▶ Maximizing the margin leads to minimizing $||W||$ i.e. None of the points fall inside the margin

$$w^*, t^* = \arg \min_{w, t} \frac{1}{2} ||w||^2 \quad \text{subject to } y_i(w \cdot x_i - t) \geq 1, 1 \leq i \leq n$$

$$\begin{aligned}
\Lambda(\mathbf{w}, t, \alpha_1, \dots, \alpha_n) &= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i (y_i (\mathbf{w} \cdot \mathbf{x}_i - t) - 1) \\
&= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^n \alpha_i y_i (\mathbf{w} \cdot \mathbf{x}_i) + \sum_{i=1}^n \alpha_i y_i t + \sum_{i=1}^n \alpha_i \\
&= \frac{1}{2} \mathbf{w} \cdot \mathbf{w} - \mathbf{w} \cdot \left(\sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right) + t \left(\sum_{i=1}^n \alpha_i y_i \right) + \sum_{i=1}^n \alpha_i
\end{aligned}$$

By taking the partial derivative of the Lagrange function with respect to t and setting it to 0

for the optimal threshold t we have

$$\sum_{i=1}^n \alpha_i y_i = 0$$

Similarly taking the partial derivative of the Lagrange function with respect to \mathbf{w} Lagrange multipliers define the weight vector as a linear combination of the training examples:

$$\frac{\partial}{\partial \mathbf{w}} \Lambda(\mathbf{w}, t, \alpha_1, \dots, \alpha_n) = \frac{\partial}{\partial \mathbf{w}} \frac{1}{2} \mathbf{w} \cdot \mathbf{w} - \frac{\partial}{\partial \mathbf{w}} \mathbf{w} \cdot \left(\sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right) = \mathbf{w} - \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

- ▶ Putting the values in Lagrangian we get

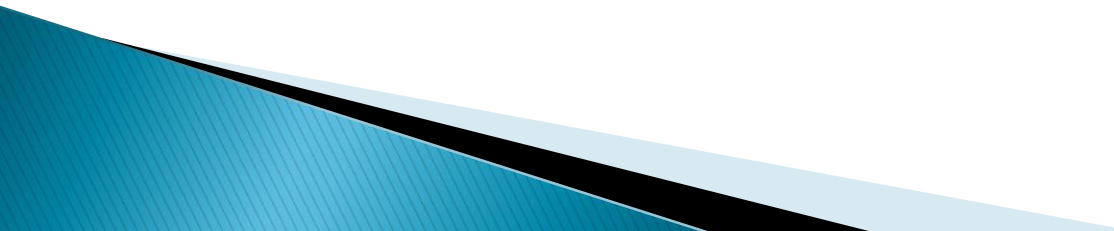
$$\mathbf{w} = \sum_{i=1}^l a_i y_i \mathbf{x}_i \quad \sum_{i=1}^n \alpha_i y_i = 0$$

$$\begin{aligned} \Lambda(\alpha_1, \dots, \alpha_n) &= -\frac{1}{2} \left(\sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right) \cdot \left(\sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right) + \sum_{i=1}^n \alpha_i \\ &= -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i \end{aligned}$$

- Maximize

$$\alpha_1^*, \dots, \alpha_n^* = \arg \max_{\alpha_1, \dots, \alpha_n} -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i$$

$$\text{subject to } \alpha_i \geq 0, 1 \leq i \leq n \text{ and } \sum_{i=1}^n \alpha_i y_i = 0$$

- ▶ Dual form of the optimization problem for SVM illustrates
 - Searching for the maximum margin decision boundary is equivalent to searching for support vectors
 - Optimization problem is defined by pair wise dot products between training examples
- 

Soft Margin SVM

- ▶ If data is not linearly separable , slack variable is introduced (one for each example)
- ▶ Which allows some of the instances to be inside margin or even at wrong side.

$$w^*, t^*, \xi_i^* = \arg \min_{w, t, \xi_i} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

subject to $y_i(w \cdot x_i - t) \geq 1 - \xi_i$ and $\xi_i \geq 0, 1 \leq i \leq n$

- ▶ *C is a user-defined parameter trading off margin maximization against slack variable*
- ▶ minimization: a high value of *C* means that margin errors incur a high penalty,
- ▶ a low value permits more margin errors in order to achieve a large margin.

If we allow more margin errors we need fewer support

- ▶ vectors, hence *C* controls to some extent the ‘complexity’ of the SVM and hence is often
- ▶ referred to as the ***complexity parameter***.

$$\begin{aligned}
\Lambda(\mathbf{w}, t, \xi_i, \alpha_i, \beta_i) &= \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i - \sum_{i=1}^n \alpha_i (y_i (\mathbf{w} \cdot \mathbf{x}_i - t) - (1 - \xi_i)) - \sum_{i=1}^n \beta_i \xi_i \\
&= \frac{1}{2} \mathbf{w} \cdot \mathbf{w} - \mathbf{w} \cdot \left(\sum_{i=1}^n \alpha_i y_i \mathbf{x}_i \right) + t \left(\sum_{i=1}^n \alpha_i y_i \right) + \sum_{i=1}^n \alpha_i + \sum_{i=1}^n (C - \alpha_i - \beta_i) \xi_i \\
&= \Lambda(\mathbf{w}, t, \alpha_i) + \sum_{i=1}^n (C - \alpha_i - \beta_i) \xi_i
\end{aligned}$$

$$\alpha_1^*, \dots, \alpha_n^* = \operatorname{argmax}_{\alpha_1, \dots, \alpha_n} -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}_i \cdot \mathbf{x}_j + \sum_{i=1}^n \alpha_i$$

$$\text{subject to } 0 \leq \alpha_i \leq C \text{ and } \sum_{i=1}^n \alpha_i y_i = 0$$

Kernel methods for non-Linearity


- ▶ Transform the data non-linearly to feature space in which linear classification can be applied.
- ▶ Training data

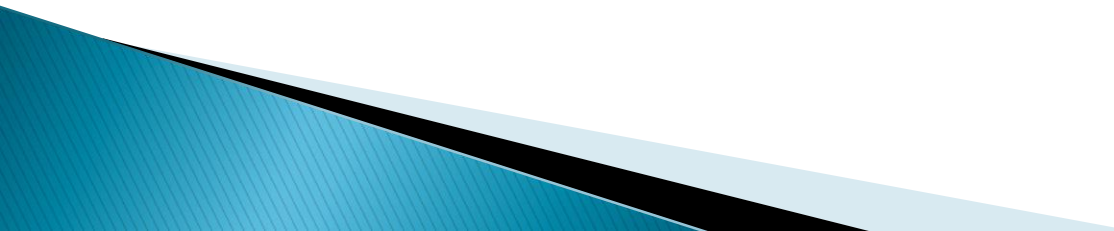



feature space






Model



- ▶ move away from a 2d view of the data to a 3d view
 - ▶ i.e. the mapping of data into a higher dimension. This is known as kernelling.
 - ▶ These are functions which takes low dimensional input space and transform it to a higher dimensional space i.e. it converts not separable problem to separable problem, these functions are called kernels.
 - ▶ It is mostly useful in non-linear separation problem.
- 





- ▶ A function that calculates the dot product in feature space directly from the vectors in the original space is called a *kernel*

$$\kappa(x1, x2) = (x1 \cdot x2)^2$$

Dot product :

$$x_i \bullet x_j = x_i x_j + y_i y_j$$

Corresponding features in quadratic space

$$(x_i^2, y_i^2) \bullet (x_j^2, y_j^2) = x_i^2 x_j^2 + y_i^2 y_j^2$$

- ▶ The instance (x,y) can be transformed into 3D feature space as $(x, y, \sqrt{2xy})$

$$\phi(\mathbf{x}_i) = (x_i^2, y_i^2, \sqrt{2}x_i y_i) \quad \phi(\mathbf{x}_j) = (x_j^2, y_j^2, \sqrt{2}x_j y_j)$$

$$\phi(\mathbf{x}_i) \cdot \phi(\mathbf{x}_j) = x_i^2 x_j^2 + y_i^2 y_j^2 + 2x_i x_j y_i y_j = (\mathbf{x}_i \cdot \mathbf{x}_j)^2$$

Kernel Method – Polynomial

- ▶ a polynomial kernel of any degree p as
 $\kappa(x_i, x_j) = (x_i \cdot x_j)^p$.
- ▶ *This transforms a d -dimensional input space into a high-dimensional feature space, such that each new feature is a product of p terms*

Kernel Method–Polynomial

- Polynomials of degree d
- Polynomials of degree up to d

$$K(\mathbf{u}, \mathbf{v}) = (\mathbf{u} \cdot \mathbf{v})^d$$

all lower-order terms are also included

$$K(\mathbf{u}, \mathbf{v}) = (\mathbf{u} \cdot \mathbf{v} + 1)^d$$

$$\phi(\mathbf{x}) = (x^2, y^2, \sqrt{2}xy, \sqrt{2}x, \sqrt{2}y, 1)$$

Kernel Method– Gaussian

- ▶ Gaussian Method

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(\frac{-\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$

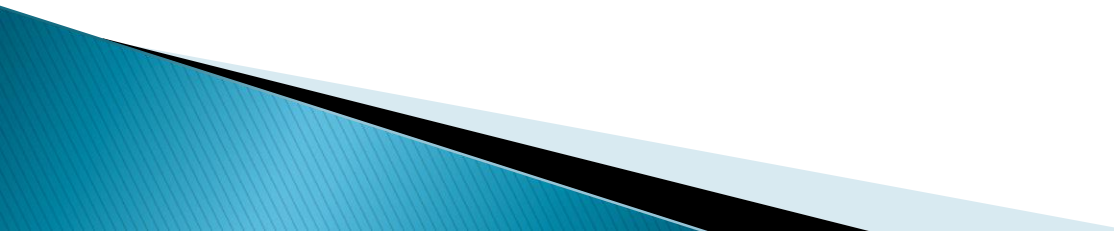
Where σ is known as bandwidth

- ▶ Kernel methods with support vector machines

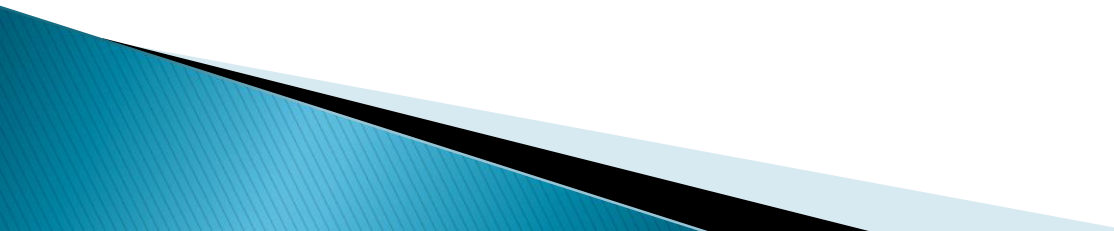
$$\alpha_1^*, \dots, \alpha_n^* = \arg \max_{\alpha_1, \dots, \alpha_n} -\frac{1}{2} \sum_{l=1}^n \sum_{j=1}^n \alpha_l \alpha_j y_l y_j \kappa(\mathbf{x}_l, \mathbf{x}_j) + \sum_{l=1}^n \alpha_l$$

$$\text{subject to } 0 \leq \alpha_l \leq C \text{ and } \sum_{l=1}^n \alpha_l y_l = 0$$


Hard Margin SVM vs. Soft margin SVM

1. Hard margin given by Boser et al. 1992 in COLT and soft margin given by Vapnik et al. 1995.
 2. Soft margin is extended version of hard margin SVM.
 3. Hard margin SVM can work only when data is completely linearly separable without any errors (noise or outliers). In case of errors either the margin is smaller or hard margin SVM fails.
 4. On the other hand soft margin SVM was proposed by Vapnik to solve this problem by introducing slack variables.
 5. As for as their usage is concerned since Soft margin is extended version of hard margin SVM so we use Soft margin SVM.
- 


Soft Margin

- ▶ try to find a line to separate, but tolerate one or few misclassified instances
 - ▶ Two types of misclassifications are tolerated by SVM under soft margin:
 - ▶ The instance is on the wrong side of the decision boundary but on the correct side/ on the margin
 - ▶ The instance is on the wrong side of the decision boundary and on the wrong side of the margin
- 

Linear Not Separable 1



Linear Not Separable 2



▶ **Pros**

- Accuracy
- Works well on smaller cleaner datasets
- It can be more efficient because it uses a subset of training points

▶ **Cons**

- Isn't suited to larger datasets as the training time with SVMs can be high
 - Less effective on noisier datasets with overlapping classes
- 